

# ML For Business II Final Project

## Predicting YouTube Spam Comments

### Group Members:

Elie LAFORGE

Yunan WANG

## Introduction

- Objective: Predict whether a YouTube comment is spam ( 1 ) or not ( 0 ).
- Dataset sourced from Kaggle, labeled in the `CLASS` column.

## Dataset Overview

- Dataset contains labeled comments from YouTube.
- Key features include comment\_id, author, date, content and video name.
- Target column: `CLASS` (1 = spam, 0 = not spam).

# Exploratory Data Analysis (EDA)

## Label Distribution

- Analyze distribution of spam vs. non-spam comments.

# Exploratory Data Analysis (EDA)

# Word Cloud

- Visualizing common words in spam comments



# Data Preprocessing

- Data cleaning steps:
  - Handling missing values.
  - Removing duplicates.
  - Text vectorization using TF-IDF.

## Model Selection

- Tried different models:
  - Logistic Regression (tested with both TF-IDF and Count Vectorizer)
  - Bernoulli Naive Bayes
  - XGBoost

# Model Evaluation

- Metrics used:
  - Accuracy
  - Precision
  - Recall
  - F1 Score
  - ROC-AUC (Receiver Operating Characteristic Area Under the Curve)
- Compared performance across models.

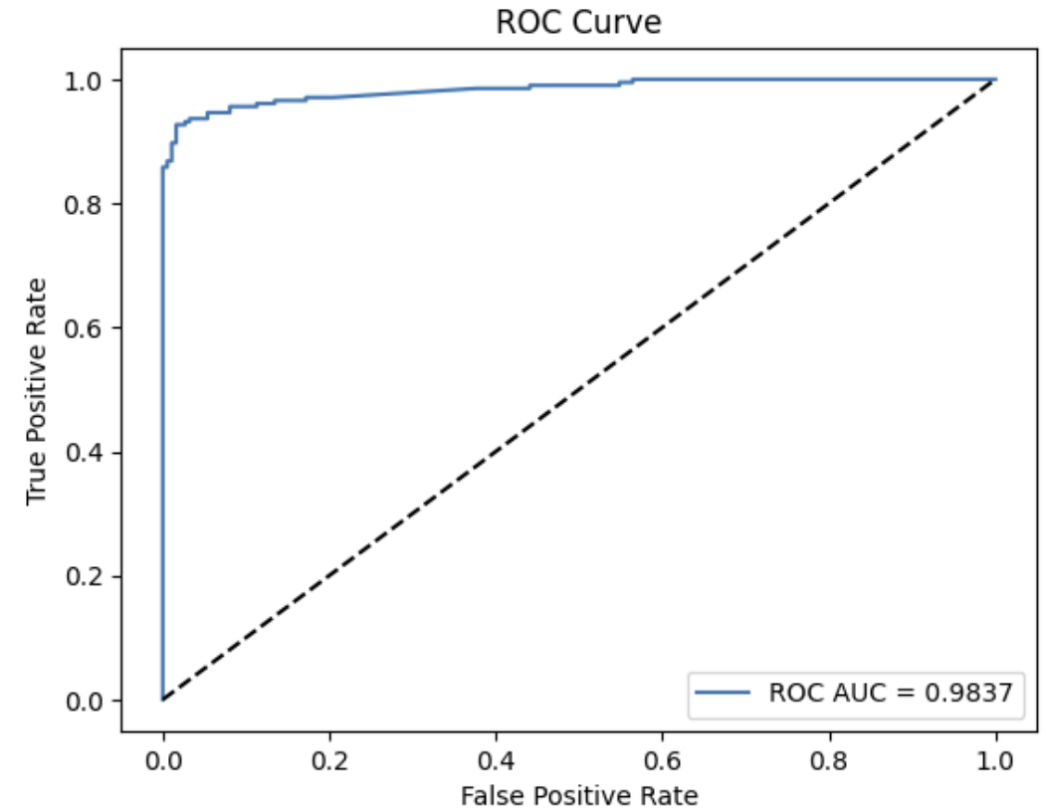


## Models tried

Metrics	Log Reg	Naive Bayes	XGBoost
Accuracy	0.90	0.78	0.90
F1-Score	0.89	0.73	0.89
ROC-AUC	0.95	0.90	0.96

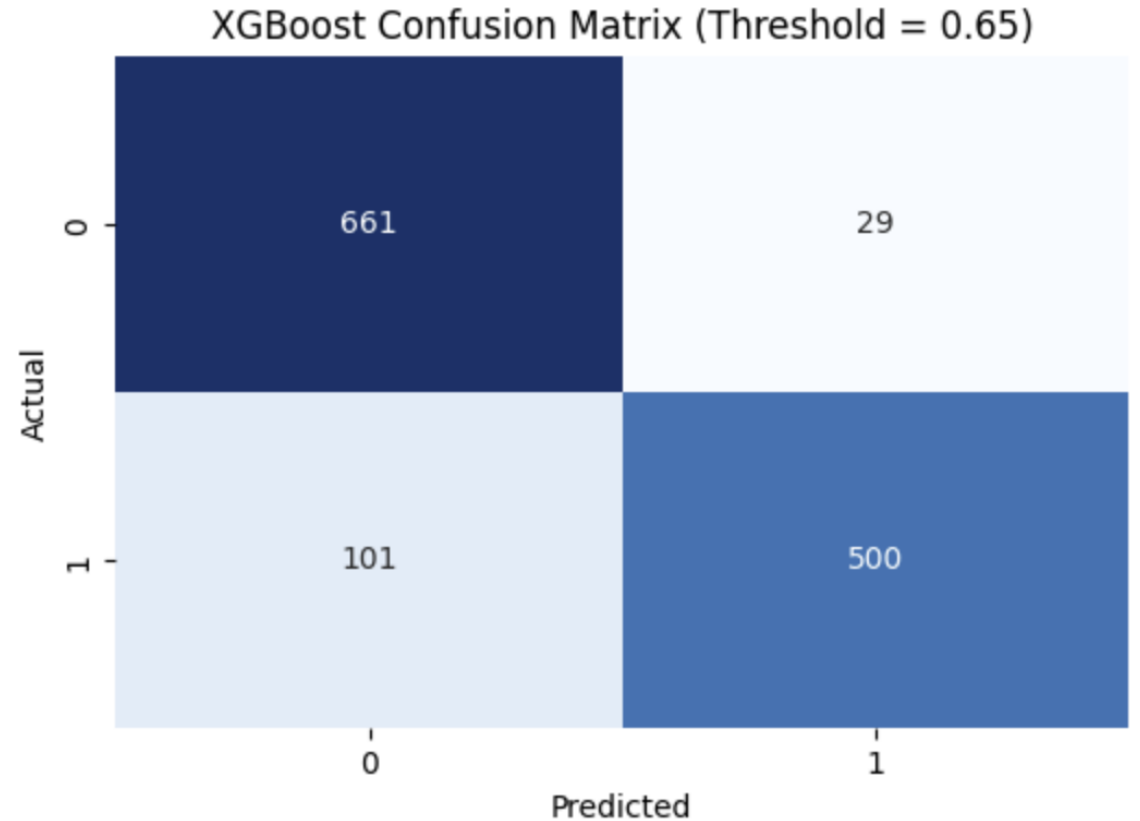
## Best Model : XGBoost

- Best performance observed with XGBoost.
- TF-IDF vectorization + hyperparameter tuning.
- Performance metrics:
  - Accuracy on test: 0.95
  - F1 Score: 0.95
  - ROC-AUC : 0.98



## Results Visualization

- Confusion matrix of the best model.
- Hyperparameters tuned:
  - learning rate = 0.2
  - max depth = 6
  - n estimators = 150



## Recap

- XGBoost + TF-IDF gave the best performance.
- Spam detection accuracy: 90%.

## DevOps Structure: Chrome Extension

Frontend PoC with Chrome Extension

- Connect to backend API server deployed on Heroku.
- Dynamically insert a pill into the comment section to indicate the spam probability.

## DevOps Structure: Testing

Unit testing with pytest on API server, covering all the endpoints and edge cases; For example:

- requests with empty content;
- malformed requests;
- internal server errors (model loading failure, etc.)

# DevOps Structure: CI/CD

## Main tool: GitHub Actions

- Code quality check on commits & PRs (format, linting, unit testing)
- Build Docker image and push to GitHub Container Registry
- Deploy to Heroku worker
- Compile this Marp deck and publish to GitHub Pages

### All workflows

Showing runs from all workflows

Q Filter workflow runs

33 workflow runs			Event ▾	Status ▾	Branch ▾	Actor ▾
✓	<b>pages build and deployment</b>					
pages-build-deployment #1: by mintyfrankie						
						17 minutes ago ...
						44s
✓	<b>github: add presentation workflow</b>					
marp-to-pages #1: Commit <a href="#">77bd715</a> pushed by mintyfrankie					main	
						21 minutes ago ...
						56s
✓	<b>github: add presentation workflow</b>					
Test #28: Commit <a href="#">77bd715</a> pushed by mintyfrankie					main	
						21 minutes ago ...
						53s
✓	<b>fix(extension): adopt incremental comme...</b>					
Test #27: Commit <a href="#">a9da95e</a> pushed by mintyfrankie					main	
						yesterday ...
						53s
✓	<b>fix(dependencies): reconfirm fastapi dep...</b>					
Test #26: Commit <a href="#">7ff1601</a> pushed by mintyfrankie					main	
						2 days ago ...
						1m 7s

## DevOps Structure: Experiment Tracking

- MLFlow for experiment tracking on Databricks
  - Two environments: staging and production
  - Log parameters, metrics, and model artifacts.
- Leverage `mlflow.autolog()` to automatically log parameters, metrics, and model artifacts.



**Live demo time!**

## Next Steps

- Potential improvements:
  - Use deep learning models (e.g., LSTM, BERT).
  - Experiment with more feature engineering.
  - Use an even larger dataset to have an even more general model

**- Sell it to Youtube and live off the profits!**